

Java Programming 2

Lecture #15

Mary Ellen Foster

MaryEllen.Foster@glasgow.ac.uk

11 November 2019

Outline

Concurrent programming

Threads in Java

Concurrent programming: basics

Process

Self-contained execution environment – own memory space

Often synonymous with programs or applications – however, an application may consist of several processes (e.g., Google Chrome)

Thread

Lightweight processes: shared memory space

Every process has at least one thread



<https://www.howtogeek.com/124218/why-does-chrome-have-so-many-open-processes/>

Why use concurrent techniques?

Tasks can be divided into subtasks; subtasks can be executed in parallel

Theoretical possible performance gain (Amdahl's Law):

If F is the percentage of the program which cannot run in parallel and n is the number of processes, then the maximum performance gain is $\frac{1}{F + \left(\frac{1-F}{n}\right)}$

Why NOT use concurrent techniques?

Threads can access shared data – two main potential problems

Visibility

Thread A reads shared data which is later changed by thread B without thread A being aware of the change

Access

Several threads access and change shared data at the same time

Can lead to

Liveness failure – program does not react any more due to problems in shared access

Safety failure – program creates incorrect data

Concurrent programming in Java

Java uses **multithreaded** programming within a single process

Basic building block: **Thread** class

Useful helper package: **java.util.concurrent** (since Java 1.5)

Note: all work we have done so far is taking place in the context of a single main **Thread** object – can be accessed and manipulated just like other Threads

Creating a Thread

Preferred technique: implement the **Runnable** interface and define **run()** method

(Can also subclass Thread but not as flexible/general)

Create thread based on **Runnable** class and use **Thread.start()** to start it

```
public class HelloRunnable implements Runnable {  
  
    @Override  
    public void run() {  
        System.out.println("Hello from a thread");  
    }  
  
    public static void main(String[] args) {  
        Thread t = new Thread(new HelloRunnable());  
        t.start();  
    }  
}
```

Based on <https://docs.oracle.com/javase/tutorial/essential/concurrency/runthread.html>

Pausing a Thread

Use **Thread.sleep()** to pause execution for a specified period

- Lets other threads use system resources

- Wait for something time-dependent to finish

Duration can be specified in milliseconds or nanoseconds – not guaranteed to be precise, depends on underlying OS and might be interrupted

Thread.sleep() throws **InterruptedException** (checked) – indicates that another thread interrupted the sleep and the thread should terminate

If you do not have another thread that can interrupt, you can usually just ignore this exception

Thread.sleep() example (main thread)

```
String[] importantInfo = {  
    "Mares eat oats",  
    "Does eat oats",  
    "Little lambs eat ivy",  
    "A kid will eat ivy too" };  
  
for (String info : importantInfo) {  
    // Pause for (approximately) 4 seconds  
    Thread.sleep(4000);  
    // Print a message  
    System.out.println(info);  
}
```

Waiting for a thread to terminate

Use **Thread.join()** method

Can also throw **InterruptedException** – must be caught

You should do this at the end of any multithreaded program to be sure it terminates

Interrupting a Thread

From outside: call **Thread.interrupt()** on the **Thread** object

Inside the **Thread**

If you call a method that could throw **InterruptedException** (e.g, **sleep()**, **join()**), just return after it is caught

Otherwise, periodically check **Thread.interrupted()** and return if it is true

SimpleThreads example (Java tutorial)

<https://docs.oracle.com/javase/tutorial/essential/concurrency/simple.html>

Thread interference details

Interference happens when two operations running in different threads but on the same data **interleave**

Two operations have multiple steps, and the steps overlap

Seemingly simple statements can translate to multiple steps in the virtual machine:

`i++` turns into:

1. Retrieve the current value of `i`
2. Increment the retrieved value by 1
3. Store the incremented value back into `i`

Example

```
public class Counter {  
    private int c = 0;  
  
    public void increment() {  
        c++;  
    }  
}
```

```
    public void decrement() {  
        c--;  
    }  
  
    public int value() {  
        return c;  
    }  
}
```

```
}
```


Thread interference

Possible sequence with two threads both accessing memory:

1. Thread A: Retrieve `i`
2. Thread B: Retrieve `i`.
3. Thread A: Increment retrieved value; result is 1.
4. Thread B: Decrement retrieved value; result is -1.
5. Thread A: Store result in `i`; `i` is now 1.
6. Thread B: Store result in `i`; `i` is now -1.

Avoiding interference: impose an ordering

Establish a **happens-before** relationship between two statements

Actions that create **happens-before**:

- Every statement before a **Thread.start()** happens before every statement executed by that thread

- When a thread terminates and causes **Thread.join()** to return, every statement in the terminated thread happens before every statement following the join

```
int counter = 0;
counter++;
System.out.println
    (counter);
```

Synchronized methods

Additional keyword:

synchronized

Add to method header

Ensures that:

Two calls to **synchronized** methods on **the same object** cannot interleave

When a synchronized method exits, it **happens-before** any other **synchronized** method calls **on the same object**

Constructors cannot be synchronized

```
public class SynchronizedCounter {  
    private int c = 0;  
    public synchronized void increment() {  
        c++;  
    }  
    public synchronized void decrement() {  
        c--;  
    }  
    public synchronized int value() {  
        return c;  
    }  
}
```

Under the hood: Intrinsic locks

Every Java object has a lock associated with it

A thread that needs consistent access to an object fields must **acquire** the lock before access, and **release** the lock when it is done

In between, the thread **owns** the lock – no other thread can acquire it (will block on attempt)

Note that a thread can access the same lock multiple times (**re-entrant**)

Synchronized methods make implicit use of the lock

More fine-grained option: **synchronized statements**

Synchronized statements example

```
public class MsLunch {
    private long c1 = 0;
    private long c2 = 0;
    private Object lock1 =
        new Object();
    private Object lock2 =
        new Object();

    public void inc1() {
        synchronized(lock1) {
            c1++;
        }
    }

    public void inc2() {
        synchronized(lock2) {
            c2++;
        }
    }
}
```

<https://docs.oracle.com/javase/tutorial/essential/concurrency/locksinc.html>

Atomic access

Atomic action

Effectively happens all at once – cannot stop in the middle

Reads and writes are atomic for most types (except long/double)

Increments like c++ are **not** atomic

Avoids need for synchronized code

Liveness problems

Liveness: concurrent program's ability to execute in a timely fashion

Potential problems:

Deadlock: two or more threads are blocked forever, waiting for each other

Starvation: a thread cannot gain access to a shared resource and is unable to make progress

Livelock: threads too busy responding to each other to make progress

Deadlock is by far the most common problem

Next time

Higher-level concurrent API

Immutable classes

Annotations